

Wine Quality Dataset

```
library(MASS)
library(dplyr)
library(tibble)
library(knitr)
library(readxl)
library(ROCR)
library(tidyr)
library(ggplot2)
library(statsr)
library(randomForest)
library(caret)
library(e1071)
library(pROC)
library(party)
library(rpart)
```

Problem 4:

We import the wine dataset from the csv file 'wineData.csv' and name our dataframe "wine". We are given that wine quality is a categorical variable with two classes. But presently it is a numerical variable as seen from the data summary. So we convert the variable "quality" into a factor variable.

```
wine <- read.csv("wineData.csv")
summary(wine)
```

##	fixed.acidity	volatile.acidity	citric.acid	residual.sugar
##	Min. : 3.800	Min. : 0.0800	Min. : 0.0000	Min. : 0.600
##	1st Qu.: 6.300	1st Qu.: 0.2100	1st Qu.: 0.2700	1st Qu.: 1.700
##	Median : 6.800	Median : 0.2600	Median : 0.3200	Median : 5.200
##	Mean : 6.855	Mean : 0.2782	Mean : 0.3342	Mean : 6.391
##	3rd Qu.: 7.300	3rd Qu.: 0.3200	3rd Qu.: 0.3900	3rd Qu.: 9.900
##	Max. : 14.200	Max. : 1.1000	Max. : 1.6600	Max. : 65.800
##	chlorides	free.sulfur.dioxide	total.sulfur.dioxide	
##	Min. : 0.00900	Min. : 2.00	Min. : 9.0	
##	1st Qu.: 0.03600	1st Qu.: 23.00	1st Qu.: 108.0	
##	Median : 0.04300	Median : 34.00	Median : 134.0	
##	Mean : 0.04577	Mean : 35.31	Mean : 138.4	
##	3rd Qu.: 0.05000	3rd Qu.: 46.00	3rd Qu.: 167.0	
##	Max. : 0.34600	Max. : 289.00	Max. : 440.0	
##	density	pH	sulphates	alcohol
##	Min. : 0.9871	Min. : 2.720	Min. : 0.2200	Min. : 8.00
##	1st Qu.: 0.9917	1st Qu.: 3.090	1st Qu.: 0.4100	1st Qu.: 9.50
##	Median : 0.9937	Median : 3.180	Median : 0.4700	Median : 10.40

```
## Mean      :0.9940      Mean      :3.188      Mean      :0.4898      Mean      :10.51
## 3rd Qu.:0.9961      3rd Qu.:3.280      3rd Qu.:0.5500      3rd Qu.:11.40
## Max.      :1.0390      Max.      :3.820      Max.      :1.0800      Max.      :14.20
## quality
## Min.      :0.0000
## 1st Qu.:0.0000
## Median :0.0000
## Mean      :0.2164
## 3rd Qu.:0.0000
## Max.      :1.0000
```

```
wine$quality <- factor(wine$quality)
```

```
colnames(wine)[colnames(wine)=="fixed.acidity"] <- "fixed_acidity"
colnames(wine)[colnames(wine)=="volatile.acidity"] <- "volatile_acidity"
colnames(wine)[colnames(wine)=="residual.sugar"] <- "residual_sugar"
colnames(wine)[colnames(wine)=="citric.acid"] <- "citric_acid"
colnames(wine)[colnames(wine)=="free.sulfur.dioxide"] <- "free_sulphur"
colnames(wine)[colnames(wine)=="total.sulfur.dioxide"] <- "total_sulphur"
```

Splitting Training and Test sets

Let us split our data into training and test sets. We will split it in a 70:30 ratio, 70% in the training set and 30% in the test set. We also split the training data into training and validation sets.

```
set.seed(101)
wine_index <- sample.int(n = nrow(wine), size = floor(.70*nrow(wine)),
replace = F)
winetrain <- wine[wine_index,]
winetest <- wine[-wine_index,]

winetrain_index <- sample.int(n = nrow(wine),
                             size = floor(.70*nrow(winetrain)), replace = F)
winetrtrain <- wine[winetrain_index,]
winetrval <- wine[-winetrain_index,]
```

We will first do parameter tuning with following steps:

1. First, we start by creating Model with default parameters
2. We will determine the best values of “cp” and “minsplit” parameters, to come to a more fine tuned model.
3. Evaluate the final model with chosen parameters on the test data using cross-validation.

Rpart decision tree using “gini”

We construct rpart decision tree using Gini index and find the best cp value using printcp().

```
library(rpart)

k=10
n = floor(nrow(winetrain)/k)
err.vect = rep(NA,k)

for (i in 1:k) {

  s1 = ((i-1) * n+1)
  s2 = (i*n)
  subset = s1:s2

  cvr.train = winetrain[-subset,]
  cvr.test = winetrain[subset,]

  winerpartgini <- rpart(quality ~ .-quality, data = cvr.train,
                        method = "class", parms = list(split = "gini" ))

  winepredgini_rpart <- prediction(
    predict(winerpartgini, newdata = cvr.test, type = "prob")[,2],
    cvr.test$quality)

  err.vect[i] <- performance(winepredgini_rpart, "auc")@y.values

  print(paste("AUC for fold", i, ":", err.vect[i]))

}

## [1] "AUC for fold 1 : 0.815129320210574"
## [1] "AUC for fold 2 : 0.773946260589366"
## [1] "AUC for fold 3 : 0.77221269296741"
## [1] "AUC for fold 4 : 0.657050465957122"
## [1] "AUC for fold 5 : 0.739284993949173"
## [1] "AUC for fold 6 : 0.791982323232323"
## [1] "AUC for fold 7 : 0.687401334216021"
## [1] "AUC for fold 8 : 0.676605387788359"
```

```
## [1] "AUC for fold 9 : 0.788862325242971"
## [1] "AUC for fold 10 : 0.736790966386555"

print(paste("Average AUC :", mean(err.vect[[i]])))

## [1] "Average AUC : 0.736790966386555"

printcp(winerpartgini)

##
## Classification tree:
## rpart(formula = quality ~ . - quality, data = cvr.train, method = "
class",
##      parms = list(split = "gini"))
##
## Variables actually used in tree construction:
## [1] alcohol          density          fixed_acidity    free_sulphur
## [5] pH              residual_sugar sulphates
##
## Root node error: 677/3086 = 0.21938
##
## n= 3086
##
##      CP nsplit rel error  xerror    xstd
## 1 0.039143      0  1.00000 1.00000 0.033957
## 2 0.018464      2  0.92171 0.95716 0.033421
## 3 0.014771      4  0.88479 0.96160 0.033478
## 4 0.010000      8  0.82422 0.92171 0.032957
```

The best value for rpart with gini, of cp comes out to be 0.01 and best value of minsplit is 8.

```
set.seed(11123)

# Run the default model

ginituned <- rpart(quality ~ .-quality, data = winetrtrain,
                  method = "class", parms = list(split = "gini" ),
                  control = rpart.control(minsplit = 8, cp = 0.01)
)

ginicon <- predict(ginituned, winetrval, type = "class")

confusionMatrix(ginicon, winetrval$quality)

## Confusion Matrix and Statistics
##
##      Reference
## Prediction    0    1
```

```
##           0 1844  408
##           1   90  157
##
##           Accuracy : 0.8007
##           95% CI : (0.7845, 0.8162)
##      No Information Rate : 0.7739
##      P-Value [Acc > NIR] : 0.0006352
##
##           Kappa : 0.2889
##
##  McNemar's Test P-Value : < 2.2e-16
##
##           Sensitivity : 0.9535
##           Specificity : 0.2779
##      Pos Pred Value : 0.8188
##      Neg Pred Value : 0.6356
##           Prevalence : 0.7739
##      Detection Rate : 0.7379
##  Detection Prevalence : 0.9012
##      Balanced Accuracy : 0.6157
##
##      'Positive' Class : 0
##
```

Rpart decision tree using “information gain”

We construct rpart decision tree using information gain index and find the best cp value using printcp().

```
k=10
n = floor(nrow(winetrain)/k)
err.vect = rep(NA,k)

for (i in 1:k) {

  s1 = ((i-1) * n+1)
  s2 = (i*n)
  subset = s1:s2

  cvr.train = winetrain[-subset,]
  cvr.test = winetrain[subset,]

  winerpartinfo <- rpart(quality ~ .-quality, data = cvr.train,
                        method = "class", parms = list(split = "informati
on" ))

  winepredinfo_rpart <- prediction(
```

```

predict(winerpartinfo, newdata = cvr.test,
type = "prob")[,2],
cvr.test$quality )

err.vect[i] <- performance(winepredinfo_rpart, "auc")@y.values

print(paste("AUC for fold", i, ":", err.vect[i]))
}

## [1] "AUC for fold 1 : 0.793797207598993"
## [1] "AUC for fold 2 : 0.726105711761343"
## [1] "AUC for fold 3 : 0.739475618720902"
## [1] "AUC for fold 4 : 0.655089881346438"
## [1] "AUC for fold 5 : 0.719972771278741"
## [1] "AUC for fold 6 : 0.737519425019425"
## [1] "AUC for fold 7 : 0.691271579161786"
## [1] "AUC for fold 8 : 0.68948922951571"
## [1] "AUC for fold 9 : 0.761472896419105"
## [1] "AUC for fold 10 : 0.737106092436975"

print(paste("Average AUC :", mean(err.vect[[i]])))

## [1] "Average AUC : 0.737106092436975"

printcp(winerpartinfo)

##
## Classification tree:
## rpart(formula = quality ~ . - quality, data = cvr.train, method = "
class",
##      parms = list(split = "information"))
##
## Variables actually used in tree construction:
## [1] alcohol      chlorides      density      fixed_acidity free_su
lphur
## [6] pH
##
## Root node error: 677/3086 = 0.21938
##
## n= 3086
##
##      CP nsplit rel error  xerror    xstd
## 1 0.039143      0  1.00000 1.00000 0.033957
## 2 0.018464      2  0.92171 0.96160 0.033478

```

```
## 3 0.012186      4   0.88479 0.94535 0.033268
## 4 0.010000      8   0.83604 0.93501 0.033133
```

The best value for rpart with information gain, of cp comes out to be 0.01 and best value of minsplit is 8.

```
infotuned <- rpart(quality ~ .-quality, data = winetrtrain,
                    method = "class", parms = list(split = "information" ),
                    control = rpart.control(minsplit = 8, cp = 0.01)
)

infocon <- predict(infotuned, winetrval, type = "class")

confusionMatrix(infocon, winetrval$quality)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction    0    1
##      0 1844  408
##      1   90  157
##
##              Accuracy : 0.8007
##              95% CI : (0.7845, 0.8162)
##      No Information Rate : 0.7739
##      P-Value [Acc > NIR] : 0.0006352
##
##              Kappa : 0.2889
##
##  Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.9535
##              Specificity : 0.2779
##      Pos Pred Value : 0.8188
##      Neg Pred Value : 0.6356
##              Prevalence : 0.7739
##      Detection Rate : 0.7379
##      Detection Prevalence : 0.9012
##      Balanced Accuracy : 0.6157
##
##              'Positive' Class : 0
##
```

The AUC for both Gini and information gain come to be the same. But the accuracy for Gini rpart is slightly higher than that of rpart information gain. So we will go ahead and check the performance of rpart using Gini on the test data

10-Fold Cross Validation for Model Evaluation on test data using Gini:

In a classification problem, we can measure the model performance using metrics like Area under the ROC curve or accuracy/error rate. We will use Area under the ROC curve as our measure of performance. To perform 10-fold cross validation, we have used a for loop which will run 10 times on different subsets of the “wine” data and generate a resulting Area under the ROC curve. The whole data is subsetted in 10 equal parts and each time the loop runs, one of those 10 is kept apart as the test set and remaining 9 subsets become a part of train set for that cross validation loop.

In the loop, we build the model using train and predict on test. We have used the prediction() and performance() functions, to calculate the area under curve for each loop(fold). After we get the AUC for all 10 folds, we take an average of them to find the overall performance of the model.

```
k=10
n = floor(nrow(wine)/k)
aucgini.vect = rep(NA,k)

for (i in 1:k) {

  s1 = ((i-1) * n+1)
  s2 = (i*n)
  subset = s1:s2

  cvrwine.train = wine[-subset,]
  cvrwine.test = wine[subset,]

  winerpartgini <- rpart(quality ~ .-quality, data = cvrwine.train,
                        method = "class", parms = list(split = "gini" ),
                        control = rpart.control(minsplit = 8, cp = 0.01))

  winepredgini_rpart <- prediction(
    predict(winerpartgini, newdata = cvrwine.test, type = "prob")[,2],
    cvrwine.test$quality )

  aucgini.vect[i] <- performance(winepredgini_rpart, "auc")@y.values

  print(paste("AUC for fold", i, ":", aucgini.vect[i]))

}

## [1] "AUC for fold 1 : 0.80751413950983"
## [1] "AUC for fold 2 : 0.692840223944875"
## [1] "AUC for fold 3 : 0.747063027649975"
```



```

## [1] "AUC for fold 4 : 0.800494653170927"
## [1] "AUC for fold 5 : 0.791452205882353"
## [1] "AUC for fold 6 : 0.830600970176069"
## [1] "AUC for fold 7 : 0.769527235354573"
## [1] "AUC for fold 8 : 0.757295907435183"
## [1] "AUC for fold 9 : 0.666986911723754"
## [1] "AUC for fold 10 : 0.730028488337587"

print(paste("Average AUC :", mean(aucgini.vect[[i]])))

## [1] "Average AUC : 0.730028488337587"

winetuned <- rpart(quality ~ .-quality, data = winetrain,
                  method = "class", parms = list(split = "gini" ),
                  control = rpart.control(minsplit = 8, cp = 0.01)
)

tunedcon <- predict(winetuned, winetest, type = "class")

confusionMatrix(tunedcon, winetest$quality)

## Confusion Matrix and Statistics
##
##              Reference
## Prediction      0      1
##              0 1108   221
##              1   49    92
##
##              Accuracy : 0.8163
##              95% CI : (0.7956, 0.8358)
##              No Information Rate : 0.7871
##              P-Value [Acc > NIR] : 0.002986
##
##              Kappa : 0.3146
##
##              Mcnemar's Test P-Value : < 2.2e-16
##
##              Sensitivity : 0.9576
##              Specificity : 0.2939
##              Pos Pred Value : 0.8337
##              Neg Pred Value : 0.6525
##              Prevalence : 0.7871
##              Detection Rate : 0.7537
##              Detection Prevalence : 0.9041
##              Balanced Accuracy : 0.6258
##

```

```
##          'Positive' Class : 0
##
```

Finding:

We get an AUC of 0.737 on the training data and an AUC of 0.73 when performing cross validation on the test data.

The accuracy on training data came to be 0.8203 and that on test data it is 0.8163.